

AI-First Company Guide

Building an AI-First Business on AI Agent Canvas

Applying re-engineering principles to multi-agent enterprise copilots

Concepts informed by *The AI-First Company* by Benjamin Simkin

Framework implementation guide for AI Agent Canvas

Table of Contents

Part 1: The AI-First Stance

- Why Re-Engineer, Not Retrofit
- The AI-First Question Applied to Your Business
- Context Is the Variable, Not the Model
- Compete on Outcomes, Not Tools

Part 2: AI Agent Canvas as Your AI-First Platform

- Architecture for Re-Engineering
- The Context Library: Personas, Context, and Entities
- Skills as Encoded Business Processes
- Focused Agents with Clear Boundaries

Part 3: The Seven Business Zones

- Mapping AI Agent Canvas to Each Zone
- MCP Connections as the Integration Spine
- Human-in-the-Loop to Autonomous Graduation

Part 4: The Implementation Sequence

- The Right Order: Foundation Before Features
- 90-Day Roadmap on AI Agent Canvas
- From Operator to Architect

References

Part 1: The AI-First Stance

The mindset shift that separates transformative AI adoption from expensive tooling

Why Re-Engineer, Not Retrofit

Most businesses approach AI adoption by bolting new tools onto existing processes. A content team adds a chatbot for drafts. A developer gets a code copilot. An automation handles some email routing. The organisation chart stays the same, the offer stays the same, and the workflows stay the same. Everything just runs a little faster.

In *The AI-First Company*, Benjamin Simkin draws a sharp distinction between two approaches to AI adoption. The first, which he calls retrofitting, takes an existing business and layers AI tools on top of it. The second, which he calls AI-first building, starts with a fundamentally different question: if you were building this company today, from nothing, with the tools that now exist, what would it look like? (Simkin, 2025, Ch. 1–2)

The difference is not semantic. Retrofitting preserves the existing structure and applies incremental speed improvements. Re-engineering challenges the structure itself. The retrofitter asks how AI can make each existing process faster. The re-engineer asks whether each process should exist in its current form at all.

The Retrofit Test

Simkin offers a simple diagnostic: if you could remove all AI from your business and it would still operate in basically the same shape, you are retrofitting. The AI is cosmetic, not structural. (Simkin, 2025, Ch. 1)

An AI-first business looks different at a structural level. The team is smaller but higher-skilled. The roles that remain are strategic, creative, or relationship-based. Production work has been absorbed into systems, not outsourced to tools that a person operates. The offer is designed around new economics: when delivery costs drop by 40–60%, the offer itself can change.

This guide shows how AI Agent Canvas provides the technical platform for that structural re-engineering. But the platform is only as good as the thinking behind it. Before touching any technology, you need to do the foundational work.

The AI-First Question Applied to Your Business

The central exercise Simkin recommends is disarmingly simple. Take a blank page and list every major area of your business: team, offer, delivery, marketing, sales, and operations. Under each one, describe how it currently works in specific terms. Then ask one question for each area: if I were starting this today, with the tools that exist right now, would I design it this way? (Simkin, 2025, Ch. 2)

The exercise produces three lists. Roughly 30% of your business is well-designed and would not change much on a clean sheet—leave it alone. About 50% is a legacy design that made sense when you built it but does not hold up anymore—that is your re-engineering work. And about 20% should not exist at all, processes or roles that persist through inertia. AI does not fix that last 20%. You simply eliminate it. (Simkin, 2025, Ch. 2)

Applying This to AI Agent Canvas

When you build on AI Agent Canvas, each component maps to a re-engineering decision:

Business Area	AI-First Question	AI Agent Canvas Component
Team roles	Which tasks need human judgement vs. which follow patterns?	Task mapping → agent workflows vs. human oversight points
Offer/Delivery	What would delivery look like at 40–60% lower cost?	Agents + MCP tools handle production; humans handle strategy
Marketing	What if content production cost dropped 80%?	Personas + context library + skills for content generation
Sales process	What if everything around the conversation was automated?	Workflows + entity store for prospect research and proposals
Operations	Which workflows should not exist at all?	Eliminate first, then automate remaining with scheduled tasks

The key insight: do not start by exploring what AI Agent Canvas can do. Start by deciding what your re-engineered business should look like. Then use the framework to build it.

Context Is the Variable, Not the Model

Simkin argues that approximately 80% of AI output quality comes from the context you provide, and only about 20% comes from how you structure the prompt itself. Most people have this ratio backwards. They spend their time crafting clever prompts when they should be building comprehensive context. (Simkin, 2025, Ch. 6 and 15)

He uses the analogy of a brilliant senior contractor who just started today. They have twenty years of general experience, they work at extraordinary speed, but they know nothing about your company. Hand them a task with no briefing and you get competent but generic output. Hand them the same task with your brand guidelines, target audience profile, examples of previous work, and the specific situation—and you get something that looks like it came from inside the company. (Simkin, 2025, Ch. 5)

The context window is that contractor’s desk. Everything they need to reference has to fit on it. Pile it with every document your company has ever produced and they cannot find anything. Give them only the three or four documents that matter for the current task and they do excellent work.

What This Means for AI Agent Canvas

AI Agent Canvas is fundamentally a context management system. Every major feature exists to get the right context into the model’s working memory at the right time:

- **Personas** are your brand voice documents. They define how the agent communicates, what expertise it draws on, and what boundaries it respects. Simkin’s “brand voice document” maps directly to a persona definition.
- **Context entries** with typed tags are the context library itself. Customer avatars, competitive positioning, SOPs, industry knowledge—each stored as a tagged entry that the framework injects automatically via AIContextProviders.

- **Entities** are your product and service documentation. Offer details, pricing logic, deliverables, case studies—structured so agents can reference them when generating proposals, content, or recommendations.
- **User Profiles** store customer avatar detail. Not generic demographics but specific characteristics, preferences, and history that make every interaction contextually relevant.
- **Guardrails** encode your business rules and compliance requirements. They constrain agent behaviour so the model operates within boundaries you define.

Critically, AI Agent Canvas does not dump everything into the context window at once. The `AIContextProvider` chain selectively injects the relevant context for each request. This is the curation principle Simkin emphasises: the skill is not accumulation, it is selecting exactly what the model needs to see for a given task. (Simkin, 2025, Ch. 6)

Compete on Outcomes, Not Tools

Simkin dedicates an entire chapter to what he calls the silent advantage. He observes two groups of business owners using AI. The first group broadcasts their AI adoption—posting their tech stacks, their workflows, their prompt libraries. The second group quietly rebuilds their operations, never mentions AI publicly, and competes purely on outcomes. (Simkin, 2025, Ch. 4)

His data across client businesses showed that proposals mentioning AI capabilities won at 24%, while proposals focused on outcomes and track record won at 41%—same team, same quality, same delivery process. The only difference was framing.

For teams building on AI Agent Canvas, the implication is clear. The framework is your operational infrastructure, not your marketing message. Your agents, personas, workflows, MCP connections, and skills are the engine. Your customers should experience better outcomes, faster turnaround, and higher quality—without needing to know how it works under the hood.

The competitive advantage comes from the accumulated context, the integrated architecture, the documented processes, and the team operating at a higher level—not from the tools themselves. Anyone can adopt the same framework. Nobody can replicate the business knowledge you encode into it.

Part 2: AI Agent Canvas as Your AI-First Platform

How the framework's architecture maps to AI-first principles

Architecture for Re-Engineering

Simkin categorises AI tools into five types: chat and copilots, automations, agents, AI-embedded software, and custom-built systems. He argues that custom-built systems—the fifth category—provide the highest long-term leverage because they use your context, connect to your data, follow your logic, and improve as your business evolves. (Simkin, 2025, Ch. 7)

AI Agent Canvas is a custom-build platform. It provides the scaffolding for creating systems that are specific to your business without requiring you to build the entire infrastructure from scratch. The framework handles orchestration, tool registration, context injection, security governance, and streaming. You focus on the business logic: what agents to create, what context to provide, and what tools to connect.

The Architecture Stack

The framework is structured as four layers, each mapping to a distinct concern:

Layer	Components	AI-First Purpose
SDK	Abstractions, Core, AgentData, Skills, Scheduler, Security, SystemTools, Notifications	The engine. Orchestration, context management, tool registry, governance. Never changes per use case.
Orchestrator	AiAgentCanvas.Orchestrator (Program.cs composition root)	Wires everything together. The single place where SDK, agents, and data connections meet.
Agents	Agent.HelloWorld (your custom agents)	Business-specific agent definitions. Personas, context seeds, workflows. Your re-engineered processes live here.
Data Connections	MCP.HelloWorldData, VectorStore.Sqlite (your custom connections)	External data access. MCP tools, vector stores, APIs. How agents reach your business data.

This separation matters for re-engineering. Agents define *how* the AI behaves (via personas, context, workflows, guardrails). Data connections define *what* it can do (via tools that access external systems). Multiple agents can share the same tools, and tools can be added or removed without changing agent logic.

The Context Library: Personas, Context, and Entities

Simkin recommends building a context library of eight categories, each typically 1,000 to 4,000 words, totalling 15,000 to 35,000 words. He estimates the initial investment at 40 to 60 hours, with compounding returns on every subsequent AI task. (Simkin, 2025, Ch. 15)

AI Agent Canvas provides a structured system for exactly this. Rather than maintaining loose documents that you manually paste into chat sessions, each category of context is stored as a first-class domain in the agent-data system and injected automatically.

Mapping Simkin's Context Library to AI Agent Canvas

Context Category (Simkin)	AI Agent Canvas Component	How It Works
Brand voice	Persona	System prompt with communication style, tone rules, example phrases. Injected via PersonaContextProvider.
Customer avatar	User Profile	Detailed customer characteristics. Stored per-user, injected per request.
Offer documentation	Entity	Products, services, pricing, deliverables. Structured as named entities with metadata.
Competitive positioning	Context entry (type: positioning)	Market position, differentiators. Tagged context injected by ContextProvider.
Case studies / proof	Context entry (type: case-study)	Client results, testimonials. Retrievable by type tag.
SOPs / process docs	Workflow	Step-by-step procedures. Executable by agents, not just reference material.
FAQ / objection handling	Context entry (type: faq)	Common questions and responses. Available to sales and support agents.
Industry regulatory	Guardrail + Context entry	Compliance rules as guardrails (enforced). Industry knowledge as context entries (referenced).

The critical advantage of this approach over loose documents is persistence and automation. Once you create a persona, it is applied to every interaction with that agent. Once you add a context entry, it is available to every agent that needs it. You build the context library once and it compounds across every task, every day, without manual intervention.

Per-Agent Context Isolation

AI Agent Canvas supports per-agent data directories. Each agent can have its own context library separate from shared context. This maps to Simkin's principle of focused agents: a sales agent sees sales-relevant context, a delivery agent sees delivery-relevant context, and neither is polluted with information irrelevant to its task.

```
agent-data/ shared/ # Context inherited by all agents agent/ # Shared personas, workflows, entities user/ # Shared user profiles orchestrator/ # Orchestrator-specific context sales-agent/ # Sales agent context library delivery-agent/ # Delivery agent context library
```

Skills as Encoded Business Processes

Simkin makes a distinction that many people miss: building a skill is not prompt engineering, it is process engineering. A prompt is a one-time instruction. A skill is a reusable, versioned, documented process that encodes how your business does a specific type of work. (Simkin, 2025, Ch. 23)

He describes five components of a good skill: a clear trigger description (when to use it), a defined input (what information the skill needs), step-by-step instructions (the process), output format specifications (what the result looks like), and quality criteria (how to judge the output). (Simkin, 2025, Ch. 23)

Skills in AI Agent Canvas

AI Agent Canvas implements skills as a first-class system. The Skills project provides a skill store, skill authoring tools, and MCP connection management. Skills are not prompts saved in a file—they are registered, discoverable, and executable through the agent's tool registry.

The framework also supports skill seeds via ISkillSeed. When you create a custom agent, you define the skills it needs as seed data. On first startup, these seeds are persisted to the agent-data directory, creating a documented, versionable library of business processes.

This connects to Simkin's documentation dividend concept: the act of encoding your processes into skills simultaneously creates the documentation that makes your business more valuable, more resilient, and easier to manage. (Simkin, 2025, Ch. 16)

Focused Agents with Clear Boundaries

Simkin advocates for focused agents—one agent, one job, clear boundaries. He warns against the industry fascination with elaborate multi-agent orchestration systems where a dozen agents coordinate on complex tasks. In his experience, focused agents that do one thing well are more reliable, more maintainable, and produce better results. (Simkin, 2025, Ch. 7)

How AI Agent Canvas Supports This

AI Agent Canvas provides the infrastructure for both focused agents and controlled collaboration between them:

- **Per-agent personas:** Each agent has its own system prompt, expertise domain, and behavioural instructions. A financial analyst agent and a customer support agent are distinct entities with distinct context.
- **Per-agent data directories:** Each agent can maintain its own context library, isolated from other agents. This prevents context pollution—a key source of quality degradation.
- **IAgentMessaging:** When agents do need to collaborate, they communicate through a structured messaging interface. The current implementation uses in-process queues, but the abstraction is designed to swap to gRPC or message queues when agents need to run as independent services.
- **Handoff:** For synchronous delegation, one agent can hand off a task to another and receive the result in the same turn. The delegating agent decides what context to share and what task to assign—maintaining clear boundaries.
- **Agent Registry:** Named agents are built from persona definitions and cached. Tools like `list_available_agents` and `get_agent_info` let agents discover each other's capabilities without tight coupling.

The design principle: agents start in-process (running inside the Orchestrator) but are structured to separate into independent services later. The separation seam is `IAgentMessaging`—swap the implementation and each agent becomes its own deployable service. This gives you the simplicity of focused agents today with the option for distributed orchestration when your business demands it.

Part 3: The Seven Business Zones

Mapping AI Agent Canvas capabilities to each area of your business

Mapping AI Agent Canvas to Each Zone

Simkin identifies seven functional zones in every business, each with different leverage points for AI. The framework he recommends is to plot each zone on two axes: how much time and money the zone consumes, and how repetitive and pattern-based the work is. Zones scoring high on both axes are your first targets. (Simkin, 2025, Ch. 9)

Here is how AI Agent Canvas maps to each zone:

Zone	AI Leverage	AI Agent Canvas Implementation
1. Lead Generation	Highest for most businesses. 60–80% compression of content production.	Persona for brand voice + context library for customer avatar + skills for content types. MCP connections to ad platforms and analytics.
2. Sales	Automate around the conversation, not the conversation itself.	Workflows for proposal generation. Entity store for prospect research. Scheduled tasks for follow-up sequencing. Human handles calls.
3. Delivery	AI handles production layer. Humans handle skilled judgement.	Agent per delivery function. Skills encode quality standards. Guardrails enforce compliance. Human review at judgement points.
4. Operations	Delete first, then automate what remains.	Scheduled tasks for recurring ops. System tools for file processing. Workflows for multi-step procedures. Goals for autonomous execution.
5. Finance	Structured data, high accuracy. Never over-trust.	MCP connections to financial systems. Guardrails for accuracy requirements. Human-in-the-loop for all financial outputs.
6. People	Strong for docs and onboarding. Weak for management.	Context library for onboarding materials. Skills for job descriptions and interview prep. Not for performance management.
7. Owner	Lowest leverage. AI assists analysis; owner decides.	Agent for market research and competitor analysis. RAG over business documents. Owner makes all strategic decisions.

MCP Connections as the Integration Spine

Simkin devotes an entire chapter to MCP (Model Context Protocol), calling it the connective tissue of an AI-first business. He describes it as the standardised way for AI to connect to your business data—CRM, ad platforms, email, file storage, analytics—so that agents can work with real information rather than requiring you to copy and paste between systems. (Simkin, 2025, Ch. 25)

He also emphasises integration architecture over point tools. A collection of disconnected AI tools gives you efficiency gains. An integrated system, where tools are connected and data flows between them, gives you compounding intelligence. The difference is whether insights from one part of your business automatically inform decisions in another. (Simkin, 2025, Ch. 13 and 19)

MCP in AI Agent Canvas

AI Agent Canvas treats MCP as a first-class integration layer:

- **Dynamic MCP connections:** Connect to external MCP servers at runtime. Add a CRM connection, an analytics connection, or a document store without changing agent code.
- **MCP connection seeds:** Custom agents declare their MCP dependencies via `IMcpConnectionSeed`. The framework validates these at startup.
- **Custom data connections:** Build your own MCP tool providers (like `MCP.HelloWorldData`) that wrap any API, database, or external service as tools the agent can call.
- **Tool governance:** All MCP tools pass through the Agent Governance Toolkit before execution. The integration is secure by default.

Each MCP connection you add makes the overall system more capable. An agent that can see your CRM, your ad platform, and your email system simultaneously can surface correlations that no human analyst would find by checking three separate dashboards. This is the integration architecture Simkin describes—data flowing between systems, creating compound intelligence.

Human-in-the-Loop to Autonomous Graduation

Simkin defines three levels of human involvement: fully human, human-in-the-loop (AI proposes, human approves), and fully autonomous. His recommendation is to start everything at human-in-the-loop and graduate to autonomous only after the system has proven itself reliable over hundreds of runs, and only for processes where errors are inexpensive. (Simkin, 2025, Ch. 14)

He calls this the graduation principle: deploy with human review on every output, track how often the human changes the AI's output, and after 200–300 cycles at 95%+ approval rate, consider moving toward autonomous with spot-checking.

The Autonomy Spectrum in AI Agent Canvas

AI Agent Canvas supports the full graduation path:

Level	AI Agent Canvas Feature	When to Use
Human-in-the-loop (default)	Chat interface via CopilotKit. User sends message, agent responds, user reviews.	Everything starts here. All client-facing output. All financial decisions. All novel situations.
Scheduled with oversight	Hangfire scheduled tasks. Agent runs on schedule, results reviewed by human.	Recurring reports. Data processing. Content drafts for review. After initial validation.

Level	AI Agent Canvas Feature	When to Use
Autonomous	Goals + Work Queue + AutonomousAgentJob. Agent claims and executes work items independently.	Only after 200+ cycles at 95%+ approval rate. Only for low-stakes tasks. Only with guardrails active.

Guardrails provide the safety net at every level. Whether an agent is responding to a chat message or executing autonomously, the same policy rules constrain its behaviour. This is Simkin's stability and safety foundation applied at the framework level. (Simkin, 2025, Ch. 20)

Part 4: The Implementation Sequence

The right order for building on AI Agent Canvas

The Right Order: Foundation Before Features

Simkin is emphatic that sequencing is the most underrated factor in AI adoption. The right systems deployed in the wrong order produce waste. His data shows that building in the wrong order costs roughly three times what doing it correctly costs. (Simkin, 2025, Ch. 31)

He prescribes a specific sequence: (1) clarify business fundamentals, (2) build the context library, (3) document processes, (4) deploy on low-risk repetitive tasks, (5) connect systems, (6) build agents and custom solutions. Each step depends on the one before it.

The most common mistake is jumping to step 6. An owner sees a demo of an agent system, gets excited, spends significant money building it, and then discovers the agents produce generic output because no context library exists, or follow broken processes because no documentation exists. The technology worked. The foundation was missing. (Simkin, 2025, Ch. 31)

The AI Agent Canvas Sequence

Applied to AI Agent Canvas, the correct implementation order is:

Step 1: Business Audit (before any code)

Do Simkin's blank-sheet exercise. For each business area, answer the AI-first question. Identify what to keep, what to re-engineer, and what to eliminate. This produces your re-engineering roadmap—the target architecture for your AI-first business. No AI Agent Canvas work happens here. This is pure business thinking.

Step 2: Build the Context Library (agent-data)

Create your persona definitions, context entries, entity records, user profiles, and guardrails. Populate the agent-data/shared/ directory with the context that all agents will need. This is the 40–60 hours of foundational work. Use Simkin's recommended approach: record yourself explaining each area to an imaginary new hire, then use AI to transcribe and structure the recording.

Step 3: Document Processes as Workflows (agent-data)

Convert your key business processes into AI Agent Canvas workflow definitions. These are not just documentation—they are executable process definitions that agents can follow. Each workflow becomes a skill seed that persists on first startup.

Step 4: Deploy on Two or Three Low-Risk Tasks (first agent)

Create your first custom agent for a specific, low-risk, repetitive task. Content drafting, report generation, data summarisation—whatever scores highest on Simkin's priority map (high cost, highly repetitive). Run it human-in-the-loop. Validate that the context library produces good output before expanding.

Step 5: Connect Data Sources (MCP connections)

Add MCP connections to your business systems. Start with your most important data source—CRM, analytics platform, document store. Give the agent a task that uses the connected data to verify the integration works. Each connection makes the previous ones more valuable.

Step 6: Expand and Graduate (more agents, autonomy)

Add more specialised agents. Connect them via IAgentMessaging. Begin graduating proven tasks from human-in-the-loop to scheduled execution. Only move to autonomous execution for tasks that have demonstrated 95%+ accuracy over hundreds of cycles.

90-Day Roadmap on AI Agent Canvas

Simkin's 90-day roadmap (Simkin, 2025, Ch. 33) adapted for AI Agent Canvas:

Month 1: The Business Audit

- Complete the AI-first question exercise for every business area.
- Identify three to five high-leverage, low-risk tasks for AI.
- Clarify positioning, offer, and customer avatar in writing.
- No AI Agent Canvas development. Pure business thinking.

Month 2: The Context Library Build

- Set up the AI Agent Canvas project (clone, configure, run).
- Create persona definitions for your first agent.
- Populate context entries: brand voice, customer avatar, offer docs, SOPs.
- Define guardrails for compliance and quality boundaries.
- Create workflow definitions for your key processes.
- Test the context by chatting with the agent. Refine until output quality matches your standards.

Month 3: First Deployments

- Deploy your first agent on two to three low-risk tasks.
- Run everything human-in-the-loop. Review every output.
- Track approval rates. Note what the agent gets wrong.
- Refine context entries and persona instructions based on real usage.
- Add your first MCP connection to a business data source.
- Measure: time saved, output quality, team feedback.

At the end of 90 days, you should have a working agent producing real output for your business, with a validated context library and at least one external data connection. The system is not dramatic yet. That is how you know it is working. Simkin notes that transformation that looks dramatic at month three is usually drama, not transformation. (Simkin, 2025, Ch. 33)

From Operator to Architect

Simkin's final chapter addresses the owner's role in an AI-first business. The transition is from operator (doing the work) to architect (designing the system that does the work). An architect does not build the building. An architect designs it, specifies the materials, defines the standards, and reviews the result. (Simkin, 2025, Ch. 34)

Applied to AI Agent Canvas, the architect's job has four components:

- **Design the system:** Decide which agents to create, what context they need, how they connect to business data, and where human oversight is required.
- **Set quality standards:** Encode your standards into personas, guardrails, and workflow definitions. The system enforces your standards automatically.
- **Monitor outcomes:** Review agent outputs, track approval rates, refine context. The scheduled tasks and autonomous execution features provide the data.
- **Make strategic decisions:** Which business areas to re-engineer next, when to graduate from human-in-the-loop to autonomous, when to add new agents or data connections.

The context library, the skills, the workflows, the guardrails—these are the architect’s blueprints. They externalise your judgement into systems that operate without you being the bottleneck. AI Agent Canvas is the construction material. Your business thinking is the architecture.

Simkin observes that the identity shift is harder than the technology. Owners who built their businesses over years or decades have their competence tied to doing the work. Letting go of that, trusting documented systems to maintain quality, is emotionally difficult even when the data proves it works. (Simkin, 2025, Ch. 34)

The test: if you left for ten days, would your AI-first business run? Not perfectly, but would it run? If the answer is yes, you have made the transition from operator to architect.

References

References

Simkin, B. (2025). *The AI-First Company: Engineering a Business That Thrives in the AI Economy*. The Mastermind.

Chapter References

- Chapters 1–2: The retrofit trap and the AI-first principle — re-engineering vs. retrofitting
- Chapter 3: AI as a mirror — AI amplifies existing business quality
- Chapter 4: The silent advantage — competing on outcomes, not tools
- Chapters 5–6: How models work and the context window as the primary constraint
- Chapter 7: Five categories of AI tools — chat, automations, agents, AI software, custom-built
- Chapter 9: The seven business zones and prioritisation framework
- Chapters 13, 19: Integration architecture — connected systems vs. point tools
- Chapter 14: Human-in-the-loop vs. autonomous — the graduation principle
- Chapters 15–16: Context libraries and the documentation dividend
- Chapter 20: Stability, reliability, safety, auditability — the four foundations
- Chapter 23: Skills as encoded business processes
- Chapter 25: MCP as the connective tissue for business data
- Chapter 31: The sequencing problem — wrong order costs 3x
- Chapter 33: The 90-day re-engineering roadmap
- Chapter 34: From operator to architect

This guide applies concepts from Simkin's work to the AI Agent Canvas framework. For the complete treatment of these principles, including detailed case studies and exercises, refer to the original book.

AI Agent Canvas is an open-source framework. For technical documentation, see the AI Agent Canvas Complete Reference Guide.